



CS 115

Functional Programming

Lecture 13: May 2, 2016

The Monad Laws



Functional Programming: Spring 2016



Today

- The monad laws
- The **Maybe** monad
- Deriving the **Maybe** monad





Recap: The Monad type class

- The **Monad** type class is defined as:

```
class Monad m where
```

```
    (>>=) :: m a -> (a -> m b) -> m b
```

```
    return :: a -> m a
```

```
    (>>) :: m a -> m b -> m b
```

```
    fail :: String -> m a
```

- **Monad** is a constructor class, since **Monad** instances are type constructors (**m**)





Recap: The **Monad** type class

```
class Monad m where
```

```
    (>>=) :: m a -> (a -> m b) -> m b
```

```
    return :: a -> m a
```

```
    (>>) :: m a -> m b -> m b
```

```
    fail :: String -> m a
```

- The two fundamental **Monad** operations are **return** and **>>=**





Recap: The Monad type class

```
class Monad m where
```

```
  (>>=) :: m a -> (a -> m b) -> m b
```

```
  return :: a -> m a
```

```
  (>>) :: m a -> m b -> m b
```

```
  fail :: String -> m a
```

- **>>=** is monadic application: a monadic function (type **a -> m b**) is applied to a monadic value (type **m a**) to get a monadic value (type **m b**)
- **return** "lifts" a regular value into a monadic value
 - *i.e.* a computation "returning" that value





Recap: The Monad type class

```
class Monad m where
```

```
    (>>=) :: m a -> (a -> m b) -> m b
```

```
    return :: a -> m a
```

```
    (>>) :: m a -> m b -> m b
```

```
    fail :: String -> m a
```

- `>>` is monadic sequencing: two monadic values ("actions") are "run" one after the other in sequence
- The first monadic action normally has the type `m ()`
- The return value of `>>` is the return value of the second monadic action





Recap: The Monad type class

```
class Monad m where
```

```
    (>>=) :: m a -> (a -> m b) -> m b
```

```
    return :: a -> m a
```

```
    (>>) :: m a -> m b -> m b
```

```
    fail :: String -> m a
```

- **fail** is invoked on a pattern match failure when monadic computations are written using the **do** notation





Recap: The **Monad** type class

- As far as Haskell is concerned, any type constructor that implements the four **Monad** methods is a valid instance of the **Monad** type class
- But for a type constructor to truly "be" a monad, more is required!





The three laws of monadics

- Many interesting natural laws come in groups of three:
 - Newton's three laws of motion
 - The three laws of thermodynamics
 - Kepler's three laws of planetary motion
 - Asimov's three laws of robotics
- Monads also have three associated laws
- Of course, the "three laws of monadics" are far more important than any of those other laws 😊





The three laws of monadics

- Recall the whole point of monads:
 - to take computations with extra effects
 - and to be able to compose them as naturally as we can compose regular functions
- It's worth looking at normal function composition to see what laws *it* obeys, then see if there are any monadic versions of those laws that monadic function composition must also obey





Function composition

- Function composition is written in Haskell using the `(.)` operator and is defined as:

```
(.) :: (b -> c) -> (a -> b) -> (a -> c)
```

```
g . f = \x -> g (f x)
```

- Or if you prefer, use the `(>.>)` operator and write the arguments in a different order:

```
(>.>) :: (a -> b) -> (b -> c) -> (a -> c)
```

```
f >.> g = \x -> g (f x)
```

```
-- or: (>.>) = flip (.)
```





Identity laws

- There is an identity function **id** that takes a value and returns it unchanged, defined as:

id :: a -> a

id x = x

- What is the relationship between **id** and function composition?
- Composing an arbitrary function **f** with **id** should give...?
 - the original function **f** back!





Identity laws

- Specifically, we can define two "laws" that function composition with **id** must obey:

$$\text{id} . f = f$$

$$f . \text{id} = f$$

- In algebra, we say that **id** is a "left identity" of function composition (law 1) and a "right identity" of function composition (law 2)
- Any notion of function composition coupled with some kind of identity function should obey laws like these in order to behave in a "reasonable" way





Associativity law

- Function composition also has to be *associative*
- Consider three functions **f**, **g**, and **h**
- This must be true:

$$(f \cdot g) \cdot h = f \cdot (g \cdot h)$$

- In words: there is only one way to compose three functions **f**, **g**, and **h** together
- Which of the functions gets composed first doesn't matter; the end result is the same
- Again: this must be true for any "reasonable" notion of function composition





Laws and more laws

- Since any "reasonable" notion of function composition has to uphold three laws:
 - left identity with the identity function
 - right identity with the identity function
 - associativity
- ... we should expect that if monadic function composition is "reasonable", it should uphold three laws like this too
- In fact, this is the case





Monadic function composition

- Recall the monadic function composition operator:
 $(\Rightarrow) :: (a \rightarrow m\ b) \rightarrow (b \rightarrow m\ c) \rightarrow (a \rightarrow m\ c)$
- Defined in the **Control.Monad** module
- This is analogous to the $(>.>)$ operator for normal functions
- There is also another form with the arguments reversed:
 $(<=<) :: (b \rightarrow m\ c) \rightarrow (a \rightarrow m\ b) \rightarrow (a \rightarrow m\ c)$
- Also defined in **Control.Monad**





Monadic function composition

- Recall that monadic function composition can be defined in terms of monadic application:

```
(>=>) :: (a -> m b) -> (b -> m c) -> (a -> m c)  
f >=> g = \x -> f x >>= g
```

- The reversed form can be defined even more simply:

```
(<=<) :: (b -> m c) -> (a -> m b) -> (a -> m c)  
(<=<) = flip (>=>)
```





Monadic identity function

- The **return** function has this type signature:

return :: a -> m a

- Viewed as a function-with-effects, you could write **return**'s type signature schematically as:

[m]

return :: a -----> a

- where [m] represents the effects embodied in a particular monad
- Thus, **return** seems to be like an identity function for monads





Monad laws: the nice version

- If we consider **return** to be a monadic identity function and use the monadic composition operator **>=>**, the three monad laws have this form:
 - 1) **return** **>=>** **f** == **f**
 - 2) **f** **>=>** **return** == **f**
 - 3) (**f** **>=>** **g**) **>=>** **h** == **f** **>=>** (**g** **>=>** **h**)
- These are *identical* in form to the corresponding laws for function composition, except instead of using **(.)** and **id** we use **(>=>)** and **return**!





Monad laws: the nice version

- In words, the monad laws state that:
 1. monadic function composition is associative
 2. **return** is a left- and right-identity for monadic function composition
- These laws, if they hold, guarantee that monadic function composition "behaves normally" *i.e.* it behaves in the way you expect function composition to behave, modulo the monadic effects





Monad laws: the ugly version

- Most Haskell literature presents the monad laws not in terms of the $\gg=$ operator but in terms of the $\gg=>$ operator
- This gives rise to much less intuitive monad laws, but the translation between them is straightforward (though a bit grungy)
- Here, we simply present the ugly form and leave the derivations as an exercise





Monad laws: the ugly version

- The ugly version of the monad laws:
 1. `return x >>= f == f x`
 2. `mx >>= return == mx` -- `mx` is a monadic value
 3. `(mx >>= f) >>= g == mx >>= (\x -> (f x >>= g))`
- Advantages of the ugly version:
 1. Can use to simplify code: when you see patterns like `(return x >>= f)`, replace with just `(f x)`
 2. Can use to constrain definitions of `return` and `>>=` when defining new monads





Enforcing monad laws

- The problem with monad laws:
 - Haskell cannot enforce them!
 - (Similar to case of laws for **Functor** type class)
- Haskell is not powerful enough to use to prove theorems about whether particular instances of the type class **Monad** have definitions of **>>=** and **return** which obey the monad laws
- Haskell will even accept versions which do not obey these laws, as long as their types are correct!





Enforcing monad laws

- Therefore, it's up to the programmer who writes the **Monad** instance definition for a particular monad to make sure that the definitions of **>>=** and **return** obey the monad laws
- Usually, the definition of **>>=** follows directly from what the monad is trying to achieve
- The definition of **return** for a monad is often much less obvious
- We can use the monad laws to tell us what the "right" definition of **return** has to be





Enforcing monad laws

- We also must check that the "natural" definition of the `>>=` operator for a given monad obeys the monad laws in conjunction with the definition of `return`





Example: the **Maybe** monad

- We have already seen the **Maybe** type constructor:

```
data Maybe a =
```

```
    Nothing
```

```
  | Just a
```

- A **Maybe** type can be used as the return value of a function when that function may or may not be able to generate a value of that type
- Such functions have the general type:

```
a -> Maybe b
```





Example: the **Maybe** monad

- We have said that the purpose of monads is to represent "notions of computation" that are different from the standard notion of computation (pure functions)
- One such "notion of computation" is *"a computation that may fail"*
- Such a computation will naturally have the type **a -> Maybe b**
- Therefore, it's not unreasonable to expect that **Maybe** might be a monad





Example: the **Maybe** monad

- The purpose of the **Maybe** monad is to enable us to easily compose functions that may fail
- We'll use a trivial example:

```
f :: Integer -> Maybe Integer
```

```
f x = if x `mod` 2 == 0 then Nothing else Just (2 * x)
```

```
g :: Integer -> Maybe Integer
```

```
g x = if x `mod` 3 == 0 then Nothing else Just (3 * x)
```

```
h :: Integer -> Maybe Integer
```

```
h x = if x `mod` 5 == 0 then Nothing else Just (5 * x)
```





Example: the **Maybe** monad

- We would like to compose **f**, **g**, and **h** to get a final function **k**
- **k** will take an **Integer**, and will multiply it by **2**, then **3**, then **5** (total **30**) unless it's divisible by **2** or **3** or **5**, in which case it will return **Nothing**
- We can't use normal function composition to define **k** in terms of **f**, **g**, and **h**, because the output types of these functions (**Maybe Integer**) aren't the same as their input types (**Integer**)





Example: the **Maybe** monad

- Defining **k** in terms of **f**, **g**, and **h** is nevertheless straightforward in Haskell:

```
k :: Integer -> Maybe Integer
```

```
k x = case f x of
```

```
    Nothing -> Nothing
```

```
    Just y -> case g y of
```

```
        Nothing -> Nothing
```

```
        Just z -> h z
```





Example: the **Maybe** monad

- Problem: the code is repetitive and grungy
 - The **Nothing** \rightarrow **Nothing** line is repeated twice!
- If more functions were to be composed, the nesting would get even deeper
- We will use monads to clean this code up
- Note: unlike the case with the **IO** monad, monads are not "essential" in order to write code using the **Maybe** type constructor
- However, monads make working with **Maybe** much more convenient





Maybe: definition of $>>=$

- Let's start by defining the $>>=$ (monadic application) operator for the **Maybe** type constructor
- It will have the (specialized) type signature:
$$(>>=) :: \text{Maybe } a \rightarrow (a \rightarrow \text{Maybe } b) \rightarrow \text{Maybe } b$$
- Usually, the definition of $>>=$ can be obtained by understanding what monadic application is trying to achieve
- That will be the case here
 - (Still have to check it using the monad laws!)





Maybe: definition of $>>=$

$(>>=) :: \text{Maybe } a \rightarrow (a \rightarrow \text{Maybe } b) \rightarrow \text{Maybe } b$

$\text{Nothing } >>= f = ???$

$\text{Just } x \ >>= f = ???$

- Need to fill in the $???$ parts
- If the previous computation failed (returned **Nothing**), what would that computation composed with **f** do?
 - Fail!
 - *i.e.* it would also return **Nothing**
 - *i.e.* first equation is $\text{Nothing } >>= f = \text{Nothing}$





Maybe: definition of $>>=$

$(>>=) :: \text{Maybe } a \rightarrow (a \rightarrow \text{Maybe } b) \rightarrow \text{Maybe } b$

$\text{Nothing } >>= f = \text{Nothing}$

$\text{Just } x \ >>= f = ???$

- If the previous computation returned **Just x**, how would we "unpack" a value of type **a** to pass to **f**?
 - Just use **x**!
 - Second equation is **Just x >>= f = f x**





Maybe: definition of $>>=$

$(>>=) :: \text{Maybe } a \rightarrow (a \rightarrow \text{Maybe } b) \rightarrow \text{Maybe } b$

$\text{Nothing } >>= f = \text{Nothing}$

$\text{Just } x \ >>= f = f \ x$

- This is the "plausible" definition of $>>=$ for the **Maybe** monad
- We still have to verify it using the monad laws!
- Before that, though, we have to define **return**:

$\text{return} :: a \rightarrow \text{Maybe } a$

$\text{return } x = ???$





Maybe: definition of `return`

`return :: a -> Maybe a`

`return x = ???`

- Since the definition has to work for *any* type `a`, there aren't many choices
- The two "obvious" choices are:
 1. `return x = Nothing`
 2. `return x = Just x`





Maybe: definition of `return`

- It should seem plausible that

`return x = Nothing`

- is not the best candidate
- This does not look like an identity function!
- Let's demonstrate that this won't work, using the previous definition of `>>=` and the monad laws





Maybe: definition of `return`

- Given

`return x = Nothing`

- let's check monad law 1 (ugly form):

`return x >>= f == f x`

- Doing some substitutions:

`return x >>= f`

`= Nothing >>= f`

`= Nothing` *-- definition of >>=*

- which cannot in general be equal to `f x` for all `f`s and `x`s





Maybe: definition of `return`

- So this definition

`return x = Nothing`

- violates monad law 1, so the correct definition is:

`return x = Just x`

- We still have to check this and the definition of `>>=` against the monad laws!





Maybe: monad law 1

- Monad law 1 (ugly form):

`return x >>= f == f x`

- With our definition, we have:

`return x >>= f`

`= Just x >>= f`

`= f x -- definition of >>=`

- So monad law 1 holds





Maybe: monad law 2

- Monad law 2 (ugly form):

`mx >>= return == mx`

- `mx` can be either `Nothing` or `Just x`
- If `mx` is `Nothing`, we have:

`Nothing >>= return`

`= Nothing` -- definition of `>>=`

`= mx`





Maybe: monad law 2

- Monad law 2 (ugly form):

`mx >>= return == mx`

- `mx` can be either `Nothing` or `Just x`
- If `mx` is `Just x`, we have:

`Just x >>= return`

`= return x -- definition of >>=`

`= Just x -- definition of return`

`= mx`

- So monad law 2 holds





Maybe: monad law 3

- Monad law 3 (ugly form):

$(mx \gg= f) \gg= g == mx \gg= (\backslash x \rightarrow (f\ x \gg= g))$

- Case 1: mx is `Nothing`

$(\text{Nothing} \gg= f) \gg= g \quad \text{-- LHS}$

$= \text{Nothing} \gg= g \quad \text{-- definition of } \gg=$

$= \text{Nothing} \quad \text{-- definition of } \gg=$

$\text{Nothing} \gg= (\backslash x \rightarrow (f\ x \gg= g)) \quad \text{-- RHS}$

$= \text{Nothing} \quad \text{-- definition of } \gg=$

- OK, so case 1 checks out





Maybe: monad law 3

- Monad law 3 (ugly form):

$(mx \gg= f) \gg= g == mx \gg= (\lambda x \rightarrow (f\ x \gg= g))$

- Case 2: mx is `Just v`

`(Just v >>= f) >>= g` -- LHS

`= f v >>= g` -- definition of `>>=`

`Just v >>= (\x -> (f x >>= g))` -- RHS

`= (\x -> (f x >>= g)) v` -- definition of `>>=`

`= f v >>= g` -- function application

- Case 2 checks out, so monad law 3 holds





Maybe: Final form

- **Maybe** instance of **Monad** type class:

instance Monad Maybe where

return x = Just x

Nothing >>= f = Nothing

Just x >>= f = f x

- We proved that this instance definition is consistent with the monad laws
- So: **Maybe** is in fact a monad!
- Monadic composition with **Maybe** monadic functions "behaves in a sensible fashion"





Maybe: Final form

- **Maybe** instance of **Monad** type class:

instance Monad Maybe where

return x = Just x

Nothing >>= f = Nothing

Just x >>= f = f x

- Interestingly, **Maybe** monad doesn't use the default definition of **fail**; instead, we have:

fail _ = Nothing

- **>>** definition is equivalent to the default





Maybe: Final form

- Let's return to our example with **f**, **g**, **h**, and **k**
- We can now define **k** monadically as:

k :: Integer -> Maybe Integer

k = f ==> g ==> h

- This is much simpler than the explicit definition with nested **case** statements!
- Monads have allowed us to remove all the "boilerplate" code dealing with **Nothing** values and focus on the overall structure of the computation





Next time

- Practical interlude: Arrays
- The list monad

